

Document Intelligence Guidelines

These guidelines establish the core responsibilities and expectations for all document intelligence systems. Every system must be accurate, transparent, and able to cite its sources. Grounded generation is not optional; it is the defining characteristic of trustworthy retrieval-augmented systems.

Accuracy

- All retrieved chunks must be verified against source documents.
- Confidence scores must be exposed for downstream debugging.
- Hallucinated facts are a critical failure.

This section is essential reading for anyone building, maintaining, or auditing a document intelligence pipeline. Deviations require documented exceptions and sign-off.

Transparency

- Every answer must include citations.
- Citation snippets must be verifiable via exact-match.
- Model reasoning should be inspectable where possible.

Transparent systems allow operators to verify every step of the generation process. Citation mechanisms must be exact and reproducible.

Performance

- Retrieval latency must remain under 100 ms for about 100 chunks.
- Batch embedding should process at least 100 chunks per call.
- Re-ingest of unchanged documents must be a no-op.

Performance targets are not aspirational. They are contractual requirements that govern user experience and operational cost.

Security

- Sensitive data must be encrypted at rest.
- Access controls must be enforced per-document.
- Audit logs must capture every query for compliance.

Security is a first-class concern, not an afterthought. All access must be logged, reviewed, and subject to quarterly audits.