

Operational Procedures

This document describes the standard operating procedures for ingestion, indexing, and quality assurance. Every operator must be familiar with these steps before performing any production ingest. Deviations require supervisor approval.

- Step 1. Identify the document source and verify its authenticity.
- Step 2. Compute a content hash using SHA-256 over the entire file.
- Step 3. Check the content hash against the document table for idempotency.
- Step 4. Extract text with page boundaries preserved using a reliable parser.
- Step 5. Run the chunker with structural-boundary preference and overlap.
- Step 6. Submit chunks for embedding in batches of up to 100.
- Step 7. Store document metadata and chunk records in a single transaction.
- Step 8. Build or refresh the ANN index after ingestion completes.
- Step 9. Run smoke tests: top-k retrieval latency under 100 milliseconds.
- Step 10. Log all ingest events including skipped, failed, and partial outcomes.